

---

## **Lab 04: Streaming Data Processing with Spark**

CSC14118 Introduction to Big Data 20KHMT1

SmallData

2023-05-22

# Contents

<b>1</b>	<b>Lab 04: Streaming Data Processing with Spark</b>	<b>1</b>
1.1	List of team members . . . . .	1
1.2	Team's result . . . . .	1
1.3	Team reflection . . . . .	1
1.4	Get Twitter tweets . . . . .	3
1.4.1	Step 1: Download data from <a href="https://huggingface.co/datasets/deberain/ChatGPT-Tweets">https://huggingface.co/datasets/deberain/ChatGPT-Tweets</a> . . . . .	3
1.4.2	Step 2: Download mongoDB in ubuntu. . . . .	3
1.4.3	Step 3: Save tweets data to mongoDB by pymongo . . . . .	3
1.5	Stream tweets to Apache Spark . . . . .	4
1.5.1	Step 1: Install Apache Spark. . . . .	4
1.5.2	Step 2: Install kafka . . . . .	5
1.5.3	Step 3: Install some Python libraries . . . . .	6
1.5.4	Step 4: Write data from mongoDB to kafka . . . . .	6
1.5.5	Step 5: Streaming data . . . . .	6
1.6	Perform sentiment analysis on tweets . . . . .	8
1.7	Visualize the analytic results . . . . .	8
1.8	References . . . . .	11

# 1 Lab 04: Streaming Data Processing with Spark

## 1.1 List of team members

ID	Full Name	Completed
20120366	Pham Phu Hoang Son	100%
20120391	Ha Xuan Truong	100%
20120393	Huynh Minh Tu	100%
20120468	Nguyen Van Hai	100%

## 1.2 Team's result

Section	Point
Get Twitter tweets	1.5
Stream tweets to Apache Spark	4
Perform sentiment analysis on tweets	2
Visualize the analytic results	2.5

## 1.3 Team reflection

### Does your journey to the deadline have any bugs? How have you overcome it?

While working on lab04, "Streaming Data Processing with Spark," I came across the following issues:

- Initially, I tried to run Streaming Data with Spark on Google Colab, but unfortunately, it didn't work as expected. As a result, our team made the decision to switch to running the Python file locally.

- Another challenge we faced was the slow performance and frequent errors during the streaming process. Whenever we ran the code, it only processed a single batch of data before encountering issues.

**What have you learned after this process?**

From the experience of working on lab04 and encountering the mentioned issues, several valuable lessons can be learned:

- **Platform compatibility:** It is important to consider the compatibility of the tools and platforms being used for the project. In this case, attempting to run Streaming Data with Spark on Google Colab resulted in compatibility issues. Therefore, it is crucial to thoroughly test and ensure that the chosen platform supports the required dependencies and functionalities.
- **Local execution as an alternative:** When facing challenges with a specific platform, having a backup plan or alternative approach is essential. In this situation, switching to running the Python file locally proved to be a viable solution. It is crucial to be adaptable and prepared to modify the execution environment to overcome obstacles and ensure progress.
- **Performance optimization:** The slow speed and frequent errors during the streaming process highlight the importance of performance optimization. Stream processing can be resource-intensive, and it is essential to analyze and fine-tune the code to improve efficiency. This may involve optimizing Spark configurations, leveraging appropriate data structures, or implementing parallel processing techniques.
- **Error handling and resilience:** When working with streaming data, it is crucial to anticipate and handle errors effectively. The fact that the code processed only a single batch of data before encountering issues suggests the need for robust error handling mechanisms. This may involve implementing error monitoring, graceful failure handling, and appropriate error recovery strategies to ensure the stability and reliability of the streaming process.
- **Continuous testing and debugging:** Streaming data processing can be complex, and continuous testing and debugging are vital. It is important to regularly validate the code, test various scenarios, and closely monitor the system's behavior. Prompt identification and resolution of bugs and issues contribute to smoother and more reliable data processing.

## 1.4 Get Twitter tweets

### 1.4.1 Step 1: Download data from

<https://huggingface.co/datasets/deberain/ChatGPT-Tweets>.

### 1.4.2 Step 2: Download mongoDB in ubuntu.

```
sudo apt-get install gnupg
```

```
curl -fsSL https://pgp.mongodb.com/server-6.0.asc | \  
sudo gpg -o /usr/share/keyrings/mongodb-server-6.0.gpg \  
--dearmor
```

```
echo "deb [ arch=amd64,arm64 signed-by=/usr/share/keyrings/mongodb-  
server-6.0.gpg ] https://repo.mongodb.org/apt/ubuntu jammy/mongodb-  
org/6.0 multiverse" | sudo tee /etc/apt/sources.list.d/mongodb-org-  
6.0.list
```

```
sudo apt-get update
```

```
sudo apt-get install -y mongodb-org
```

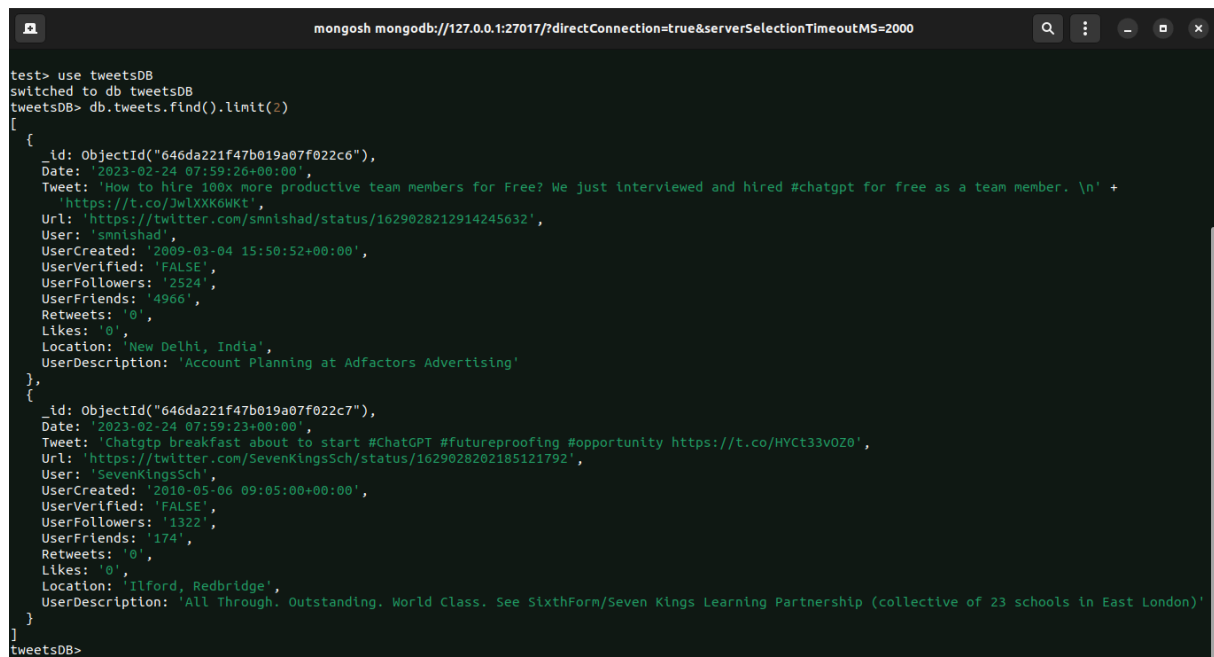
Run mongoDB in terminal:

```
sudo systemctl start mongod
```

```
sudo systemctl status mongod
```

### 1.4.3 Step 3: Save tweets data to mongoDB by pymongo

Run file data\_to\_mongo.py to save tweets data in mongoDB



```

mongosh mongodb://127.0.0.1:27017/?directConnection=true&serverSelectionTimeoutMS=2000
test> use tweetsDB
switched to db tweetsDB
tweetsDB> db.tweets.find().limit(2)
[
  {
    _id: ObjectId("646da221f47b019a07f022c6"),
    Date: '2023-02-24 07:59:26+00:00',
    Tweet: 'How to hire 100x more productive team members for free? We just interviewed and hired #chatgpt for free as a team member. \n' +
      'https://t.co/JwLXXK6Wkt',
    Url: 'https://twitter.com/snnishad/status/1629028212914245632',
    User: 'snnishad',
    UserCreated: '2009-03-04 15:50:52+00:00',
    UserVerified: 'FALSE',
    UserFollowers: '2524',
    UserFriends: '4966',
    Retweets: '0',
    Likes: '0',
    Location: 'New Delhi, India',
    UserDescription: 'Account Planning at Adfactors Advertising'
  },
  {
    _id: ObjectId("646da221f47b019a07f022c7"),
    Date: '2023-02-24 07:59:23+00:00',
    Tweet: 'Chatgpt breakfast about to start #ChatGPT #futureproofing #opportunity https://t.co/HYCT33v0Z0',
    Url: 'https://twitter.com/SevenKingsSch/status/1629028202185121792',
    User: 'SevenKingsSch',
    UserCreated: '2010-05-06 09:05:00+00:00',
    UserVerified: 'FALSE',
    UserFollowers: '1322',
    UserFriends: '174',
    Retweets: '0',
    Likes: '0',
    Location: 'Ilford, Redbridge',
    UserDescription: 'All Through. Outstanding. World Class. See SixthForm/Seven Kings Learning Partnership (collective of 23 schools in East London)'
  }
]
tweetsDB>

```

Figure 1.1: Tweets to MongoDB

## 1.5 Stream tweets to Apache Spark

We using kafka to stream tweets from mongoDB to apache Spark.

### 1.5.1 Step 1: Install Apache Spark.

Download Spark in this link: <https://www.apache.org/dyn/closer.lua/spark/spark-3.4.0/spark-3.4.0-bin-hadoop3.tgz>

Unzip and save to './Home'. Rename to Spark. Open terminal and run this command:

```
nano ~/.bashrc
```

Add and replace username to user ubuntu

```
export SPARK_HOME=/home/{username}/spark
export PATH=$PATH:$SPARK_HOME/bin:$SPARK_HOME/sbin
```

Return Terminal and run:

```
source ~/.bashrc
```

Add jars:

```
cd $SPARK_HOME/jars && wget https://repo1.maven.org/maven2/org/mongodb/spark/mongo-spark-connector_2.12/10.1.1/mongo-spark-connector_2.12-10.1.1.jar
```

```
cd $SPARK_HOME/jars && wget https://repo1.maven.org/maven2/org/mongodb/mongodb-driver/3.12.13/mongodb-driver-3.12.13.jar
```

```
cd $SPARK_HOME/jars && wget https://repo1.maven.org/maven2/org/mongodb/mongo-java-driver/3.12.13/mongo-java-driver-3.12.13.jar
```

```
cd $SPARK_HOME/jars && wget https://repo1.maven.org/maven2/org/mongodb/bson/4.9.1/bson-4.9.1.jar
```

### 1.5.2 Step 2: Install kafka

Download Kafka in this link: [https://downloads.apache.org/kafka/3.4.0/kafka\\_2.12-3.4.0.tgz](https://downloads.apache.org/kafka/3.4.0/kafka_2.12-3.4.0.tgz)

Unzip and add to '/Home'

Rename to Kafka

Open terminal and run:

```
cd /home/{username}/Kafka
```

```
bin/zookeeper-server-start.sh config/zookeeper.properties
```

Open new terminal and run:

```
cd /home/{username}/Kafka
```

```
bin/zookeeper-server-start.sh config/zookeeper.properties
```

Open new terminal and run this command to create a topic of Kafka:

```
bin/kafka-topics.sh --create --topic twitterstream --bootstrap-server localhost:9092
```

```
# verify
```

```
bin/kafka-topics.sh --describe --topic twitterstream --bootstrap-server localhost:9092
```

### 1.5.3 Step 3: Install some Python libraries

Open terminal and install some libraries

```
pip install datasets
pip install pymongo
pip install textblob
pip install findspark
```

### 1.5.4 Step 4: Write data from mongoDB to kafka

- Firstly, we convert all columns to a struct column.
- Secondly, convert struct column to JSON and rename to “value”
- Lastly, We save tweets to kafka each 15 seconds with 10000 rows

You can find source code in file `mongo_to_kafka.py` to see our step by step and get a deep understanding of how to write data to kafka.

Run file `mongo_to_kafka.py` by command:

```
spark -submit --packages org.apache.spark:spark-streaming-kafka-0-12_2:3.4.0,org.apache.spark:spark-sql-kafka-0-10_2.12:3.4.0 mongo_to_kafka.py
```

### 1.5.5 Step 5: Streaming data

ou can find source code in file `kafka_to_visualize.py` to see our step by step and get a deep understanding of how to streaming data by spark.

Run file `kafka_to_visualize.py` by command:

```
spark -submit --packages org.apache.spark:spark-streaming-kafka-0-12_2:3.4.0,org.apache.spark:spark-sql-kafka-0-10_2.12:3.4.0 kafka_to_visualize.py
```



```

electricgoal@gotcha: /media/electricgoal/Work/Big_Data/BigData_Lab04$ spark-submit --packages org.apache.spark:spark-streaming-kafka-0-10_2.12:3.4.0,org.apache.spark:spark-sql-kafka-0-10_2.12:3.4.0 kafka_to_visualize.py

use twentysDB

db.twentys.find().count()

Nếu thấy giá trị trên terminal bằng với giá trị output của file data_to_mongo thì là ok

Quay lại thư mục code và chạy lệnh sau để khởi động quá trình đọc dữ liệu từ mongo sang kafka

spark-submit --packages org.apache.spark:spark-streaming-kafka-0-10_2.12:3.4.0,org.apache.spark:spark-sql-kafka-0-10_2.12:3.4.0 mongo_to_kafka.py

Đợi một lúc, mở một terminal khác và chạy lệnh để khởi động quá trình lưu dữ liệu từ kafka sang mongo

spark-submit --packages org.apache.spark:spark-streaming-kafka-0-10_2.12:3.4.0,org.apache.spark:spark-sql-kafka-0-10_2.12:3.4.0 mongo_to_kafka.py

```

**Figure 1.2:** Command to run

```

"latestOffset" : 29,
"queryPlanning" : 59,
"triggerExecution" : 3959,
"walCommit" : 43
},
"stateOperators" : [ ],
"sources" : [ {
"description" : "KafkaV2[Subscribe[twitterstream]]",
"startOffset" : { "start" : 0 },
"twitterstream" : {
"0" : 0
}
},
"endOffset" : { "0" : 10000
},
"twitterstream" : {
"0" : 10000
}
},
"latestOffset" : {
"twitterstream" : {
"0" : 10000
}
},
"numInputRows" : 10021,
"inputRowsPerSecond" : 501.02494875256235,
"processedRowsPerSecond" : 2531.194746148017,
"metrics" : {
"avgOffsetsBehindLatest" : "0.0",
"maxOffsetsBehindLatest" : "0",
"minOffsetsBehindLatest" : "0"
}
},
"sink" : {
"description" : "ForeachBatchSink",
"numOutputRows" : -1
}
}
}

```

Dữ liệu trong mongoDB

```

{
  "data": "data",
  "type": "text",
  "value": "data"
}

```

Dữ liệu trong kafka

```

{"_id": 1, "data": "data", "type": "text", "value": "data"}
{"_id": 2, "data": "data", "type": "text", "value": "data"}
{"_id": 3, "data": "data", "type": "text", "value": "data"}
{"_id": 4, "data": "data", "type": "text", "value": "data"}
{"_id": 5, "data": "data", "type": "text", "value": "data"}
{"_id": 6, "data": "data", "type": "text", "value": "data"}
{"_id": 7, "data": "data", "type": "text", "value": "data"}
{"_id": 8, "data": "data", "type": "text", "value": "data"}
{"_id": 9, "data": "data", "type": "text", "value": "data"}
{"_id": 10, "data": "data", "type": "text", "value": "data"}

```

### Figure 1.3: Result

## 1.6 Perform sentiment analysis on tweets

The `sentiment(res_df)` function takes a DataFrame, performs text cleaning on the “Tweet” column, performs sentiment analysis, classifies the sentiments based on defined thresholds, and returns the updated DataFrame with sentiment labels and the original DataFrame for visualization purposes.

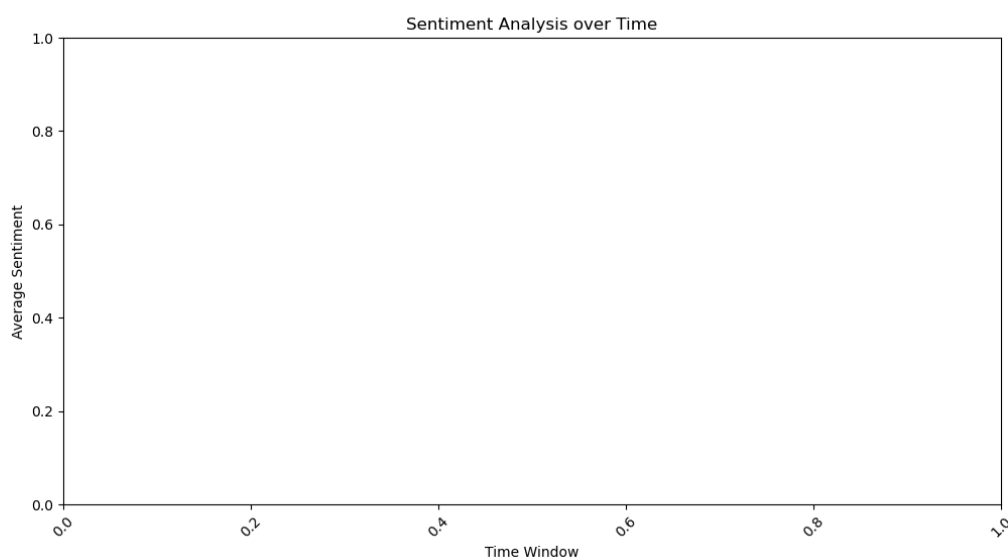
Sets the positive and negative thresholds for sentiment classification. Tweets with a sentiment score greater than `positive_threshold = 0.2` are classified as “positive”, and tweets with a sentiment score less than `negative_threshold = -0.2` are classified as “negative”. Tweets within the threshold range are classified as “neutral”.

## 1.7 Visualize the analytic results

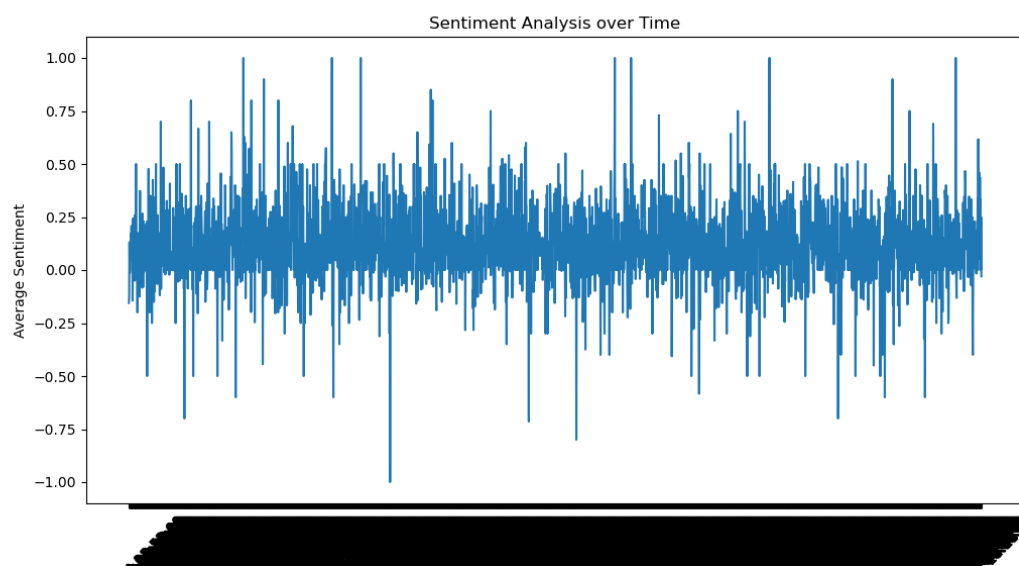
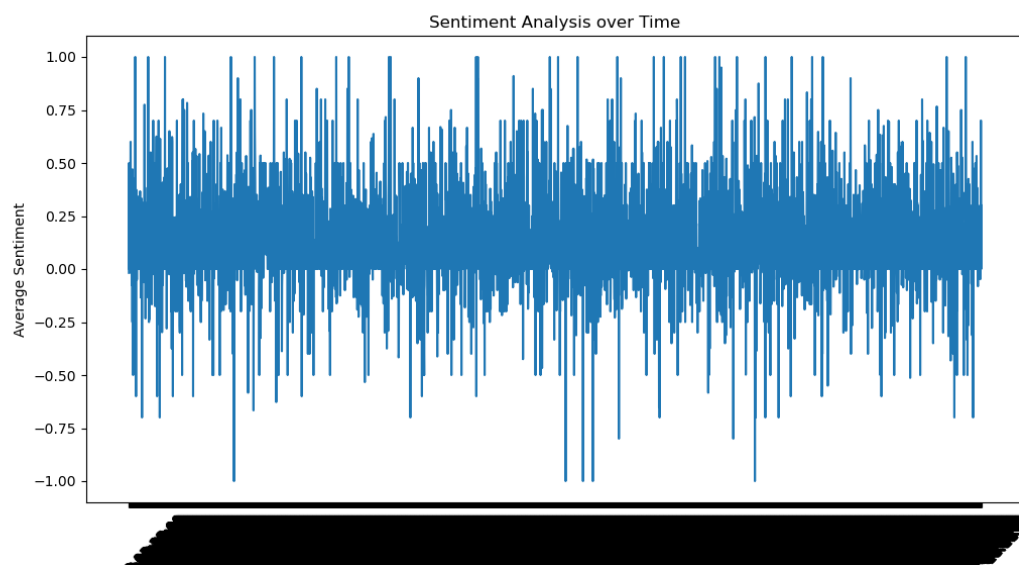
The `visualize_batch()` function takes a batch DataFrame, performs sentiment analysis on it using the sentiment function, aggregates the sentiment values over time windows, creates a line plot of the average sentiment values, and saves the plot as an image file. This function can be used iteratively for each batch of data to visualize and save sentiment analysis results over time.

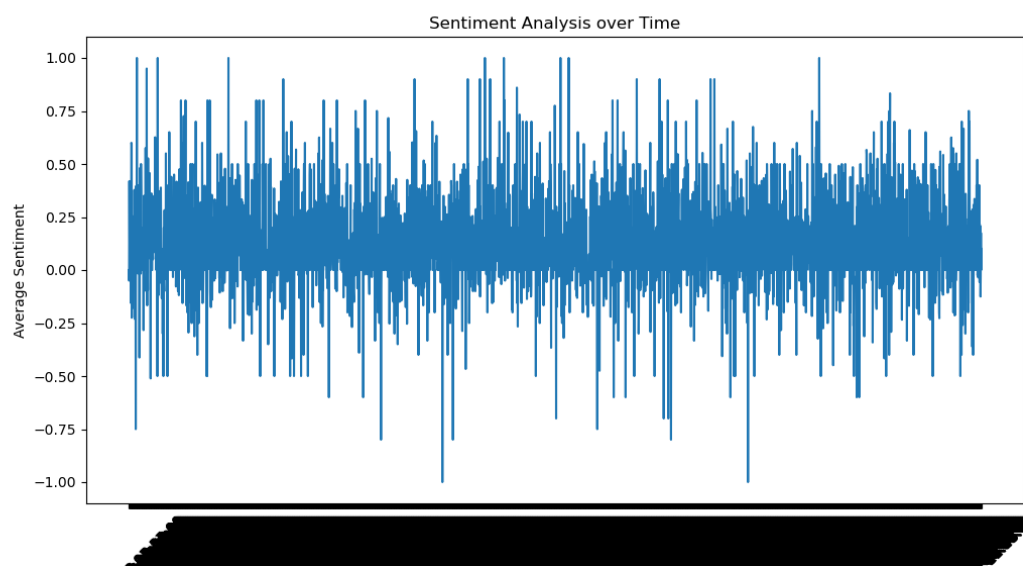
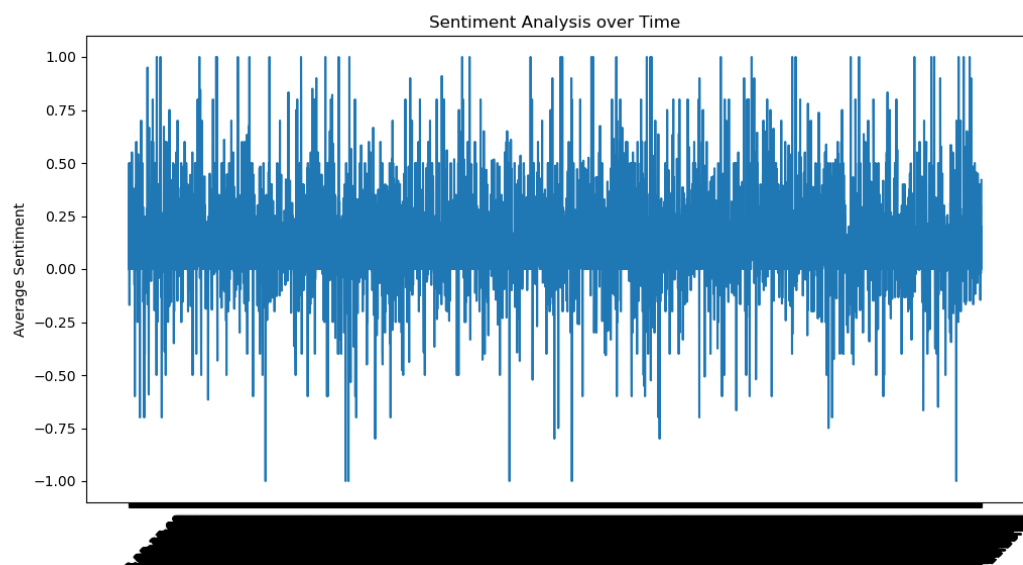
Print out the average sentiment tweets over a period (in this case is 1 minute)

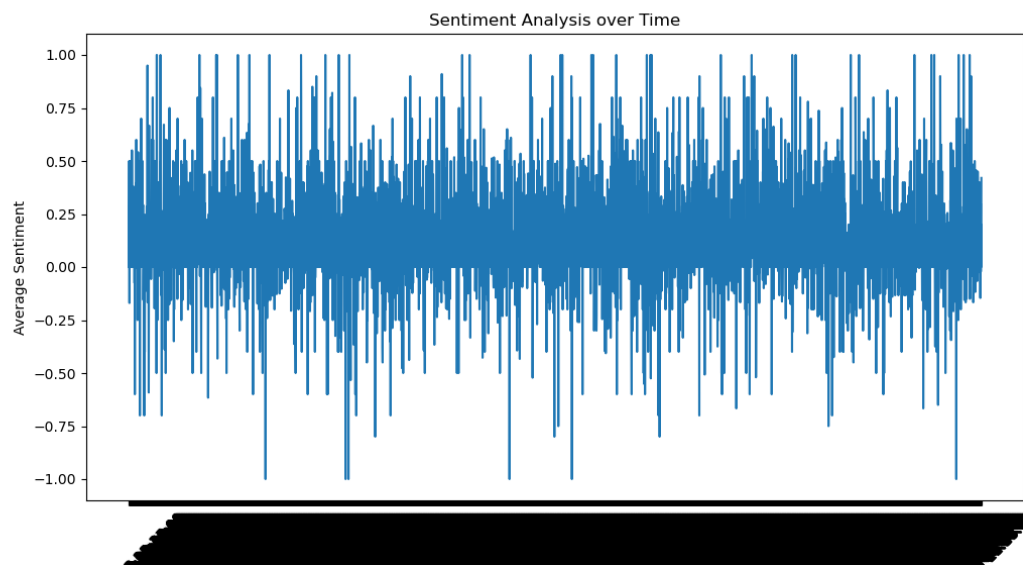
Results after running:



**Figure 1.4:** Batch 0

**Figure 1.5:** Batch 1**Figure 1.6:** Batch 2

**Figure 1.7:** Batch 3**Figure 1.8:** Batch 4



**Figure 1.9:** Batch 5

## 1.8 References

- <https://spark.apache.org/docs/latest/structured-streaming-kafka-integration.html>
- <https://kafka.apache.org/documentation/>
- [https://sparkbyexamples.com/spark/spark-streaming-with-kafka/?utm\\_content=expand\\_article](https://sparkbyexamples.com/spark/spark-streaming-with-kafka/?utm_content=expand_article)
- <https://matplotlib.org/stable/gallery/index.html>